

Pycket

A tracing JIT
For a functional language

Sam Tobin-Hochstadt
Indiana University

EPFL, 9/28/16

AOT

JIT

OO

v8, Self
HotSpot

FP

GHC, Gambit
MLton, SBCL

AOT

JIT

OO

GCJ

...

v8, Self
HotSpot

FP

GHC, Gambit
MLton, SBCL

AOT

JIT

OO

GCJ

...

v8, Self
HotSpot

FP

GHC, Gambit
MLton, SBCL

???

A simple program

```
(define (dot u v)
  (for/sum ([x u]
            [y v])
    (* x y)))
```

A simple program

```
(define (dot u v)
  (for/sum ([x u]
            [y v])
    (* x y)))
```

506 ms (size 10000000)

A simple program

```
(define (dot u v)
  (for/sum ([x (in-vector u)]
            [y (in-vector v)])
    (fl* x y)))
```

39 ms (size 10000000)

A simple program

```
(define (dot v1 v2)
  (define len (flvector-length v1))
  (unless (= len (flvector-length v2))
    (error 'fail))
  (let loop ([n 0] [sum 0.0])
    (if (unsafe-fx= len n) sum
        (loop (unsafe-fx+ n 1)
              (unsafe-fl+
               sum (unsafe-fl*
                    (unsafe-flvector-ref v1 n)
                    (unsafe-flvector-ref v1 n)))))))
```

29 ms (size 100000000)

A simple program

```
(define/contract (dot u v)
  ((vectorof flonum?) (vectorof flonum?)
   . -> . flonum?)
  (for/sum ([x (in-vector u)]
            [y (in-vector v)])
    (fl* x y)))
```

933 ms (size 10000000)

Success?

- ✓ Fast code
- ✓ Generic operations and contracts
- ✗ You can only pick one

What do contracts and
generic functions
have in common?

What do contracts and
generic functions
have in common?

Indirection

MAYBE YOU CAN
HAVE YOUR CAKE
AND EAT IT TOO.



+



PyPy

=



With added cake ...

```
(define (dot v1 v2)
  (define len (flvector-length v1))
  (unless (= len (flvector-length v2))
    (error 'fail))
  (let loop ([n 0] [sum 0.0])
    (if (unsafe-fx= len n) sum
        (loop (unsafe-fx+ n 1)
              (unsafe-fl+
               sum (unsafe-fl*
                    (unsafe-flvector-ref v1 n)
                    (unsafe-flvector-ref v1 n)))))))
```

8 ms (size 100000000)

With added cake ...

```
(define (dot u v)
  (for/sum ([x (in-vector u)]
            [y (in-vector v)])
    (fl* x y)))
```

11 ms (size 10000000)

With added cake ...

```
(define (dot u v)
  (for/sum ([x u]
            [y v])
    (* x y)))
```

12 ms (size 100000000)

With added cake ...

```
(define/contract (dot u v)
  ((vectorof flonum?) (vectorof flonum?)
   . -> . flonum?)
  (for/sum ([x (in-vector u)]
            [y (in-vector v)])
    (fl* x y)))
```

17 ms (size 10000000)

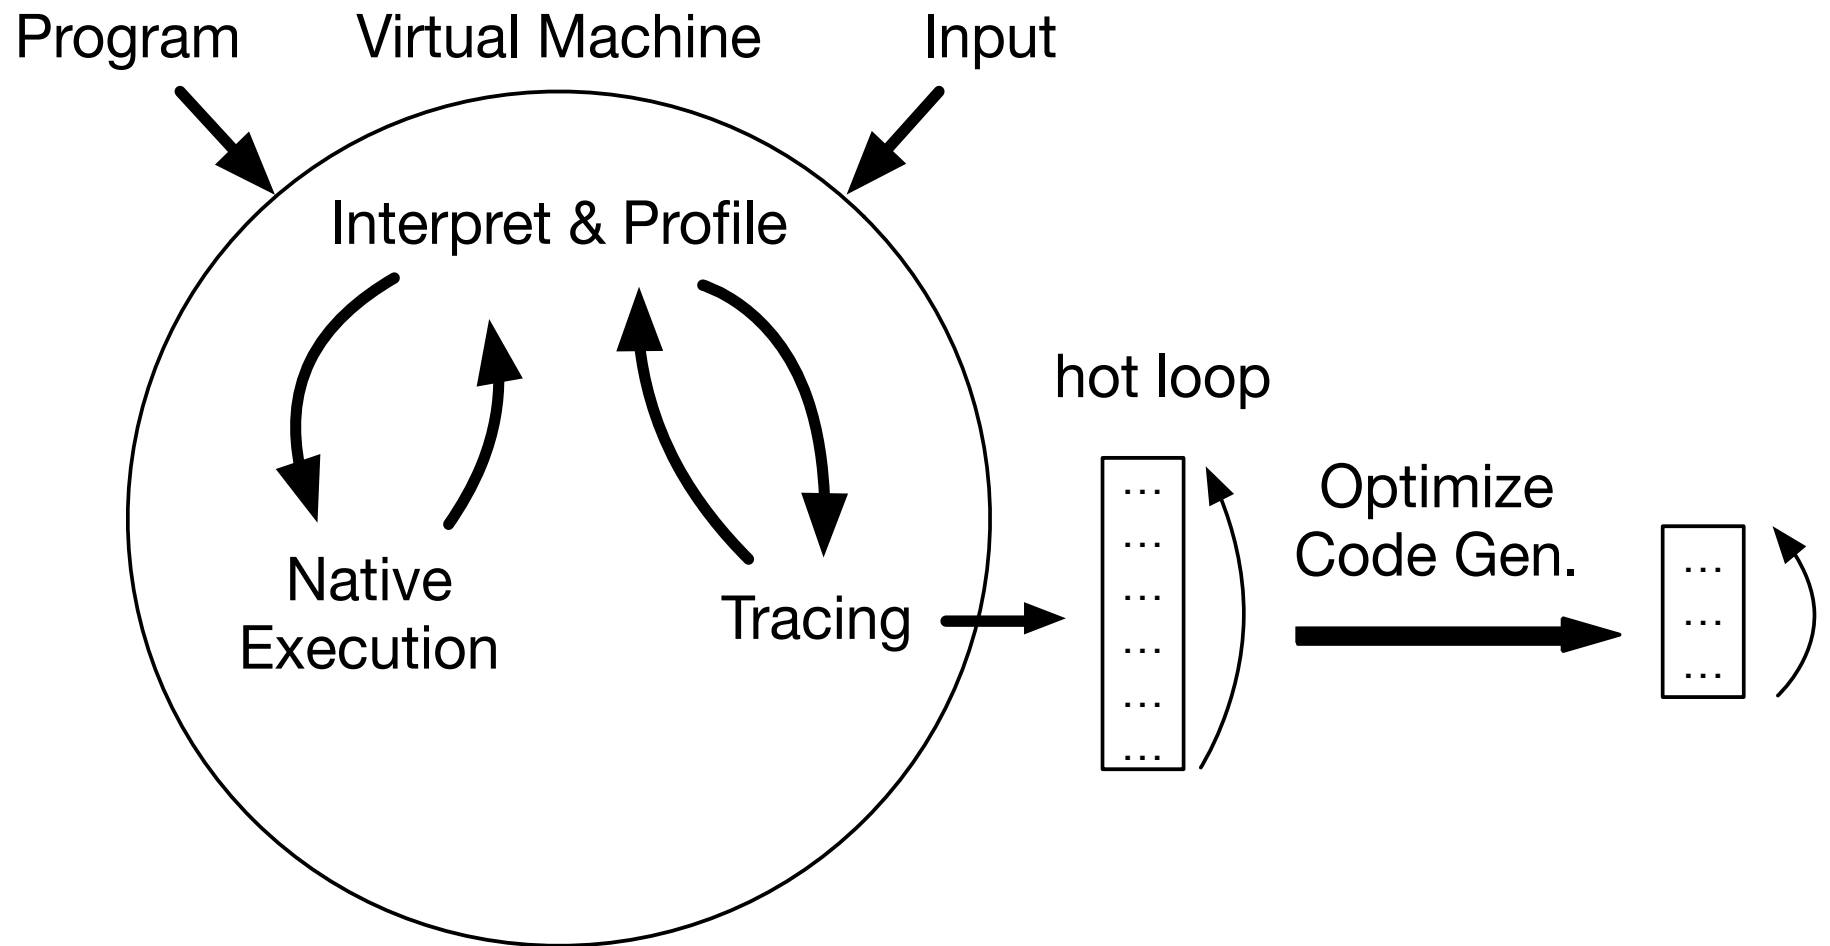


How does it work?



Tracing JIT

1. Interpret Program
2. Find hot loop
3. Record operations for one iteration
4. Optimize
5. Switch to new code

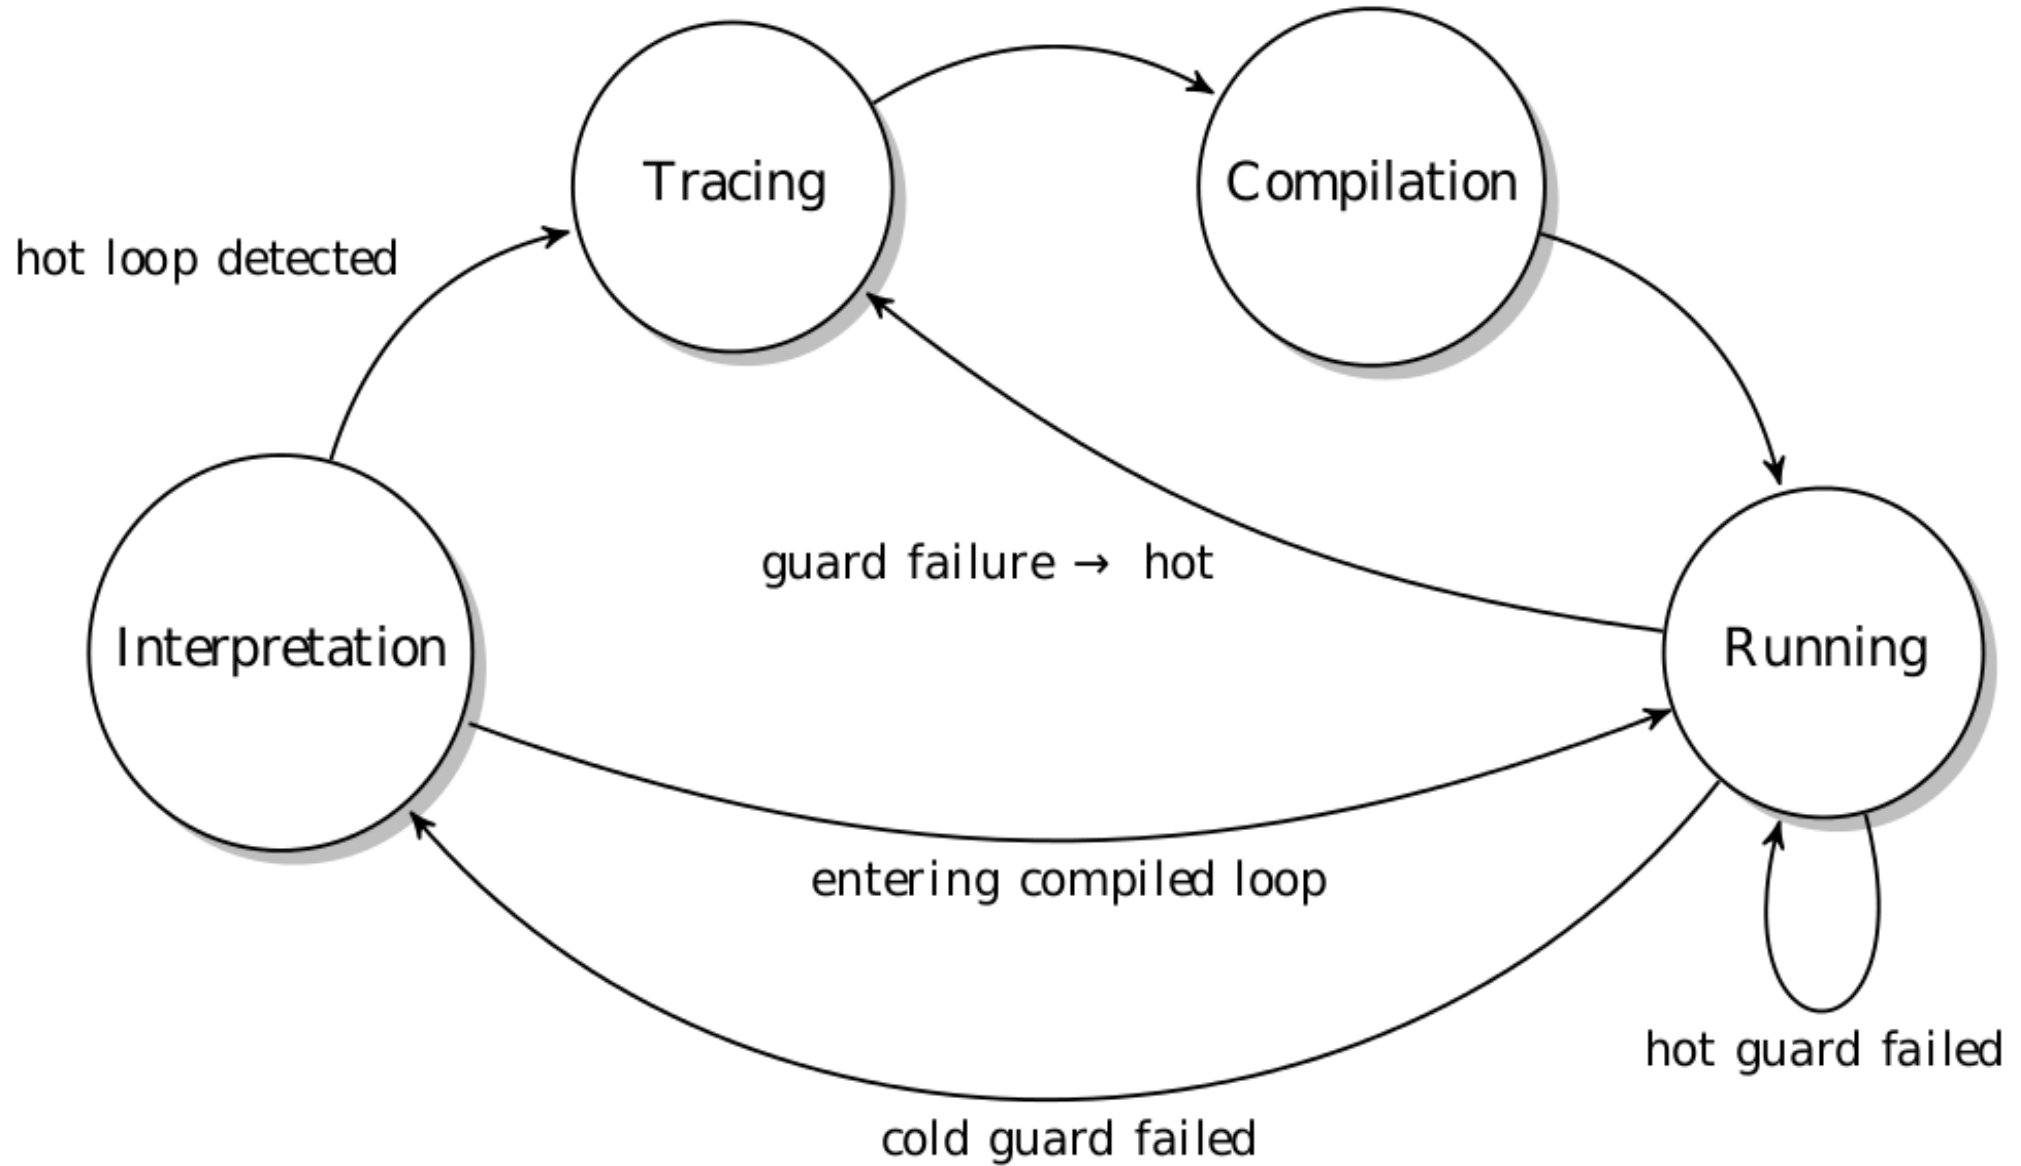
Tracing JIT



Tracing JIT



Tracing JIT



(Diagram from Antonio Cuni)

Resulting Optimizations


Inlining (happens for free)

Constant propagation

Allocation Removal

Dot product Inner Loop

```
label(acc, idx1, idx2, len1, len2, arr1, arr2)
  check loop counters
  guard(idx1 < len1)
  guard(idx2 < len2)
  fetch elements
  val1      = getarrayitem_gc(arr1, idx1)
  val2      = getarrayitem_gc(arr2, idx2)
  computation
  prod      = val1 * val2
  acc_new   = acc + prod
  increment counters
  idx1_new  = idx1 + 1
  idx2_new  = idx2 + 1
  loop back
  jump(acc_new, idx1_new, idx2_new, len1, len2, arr1, arr2)
```



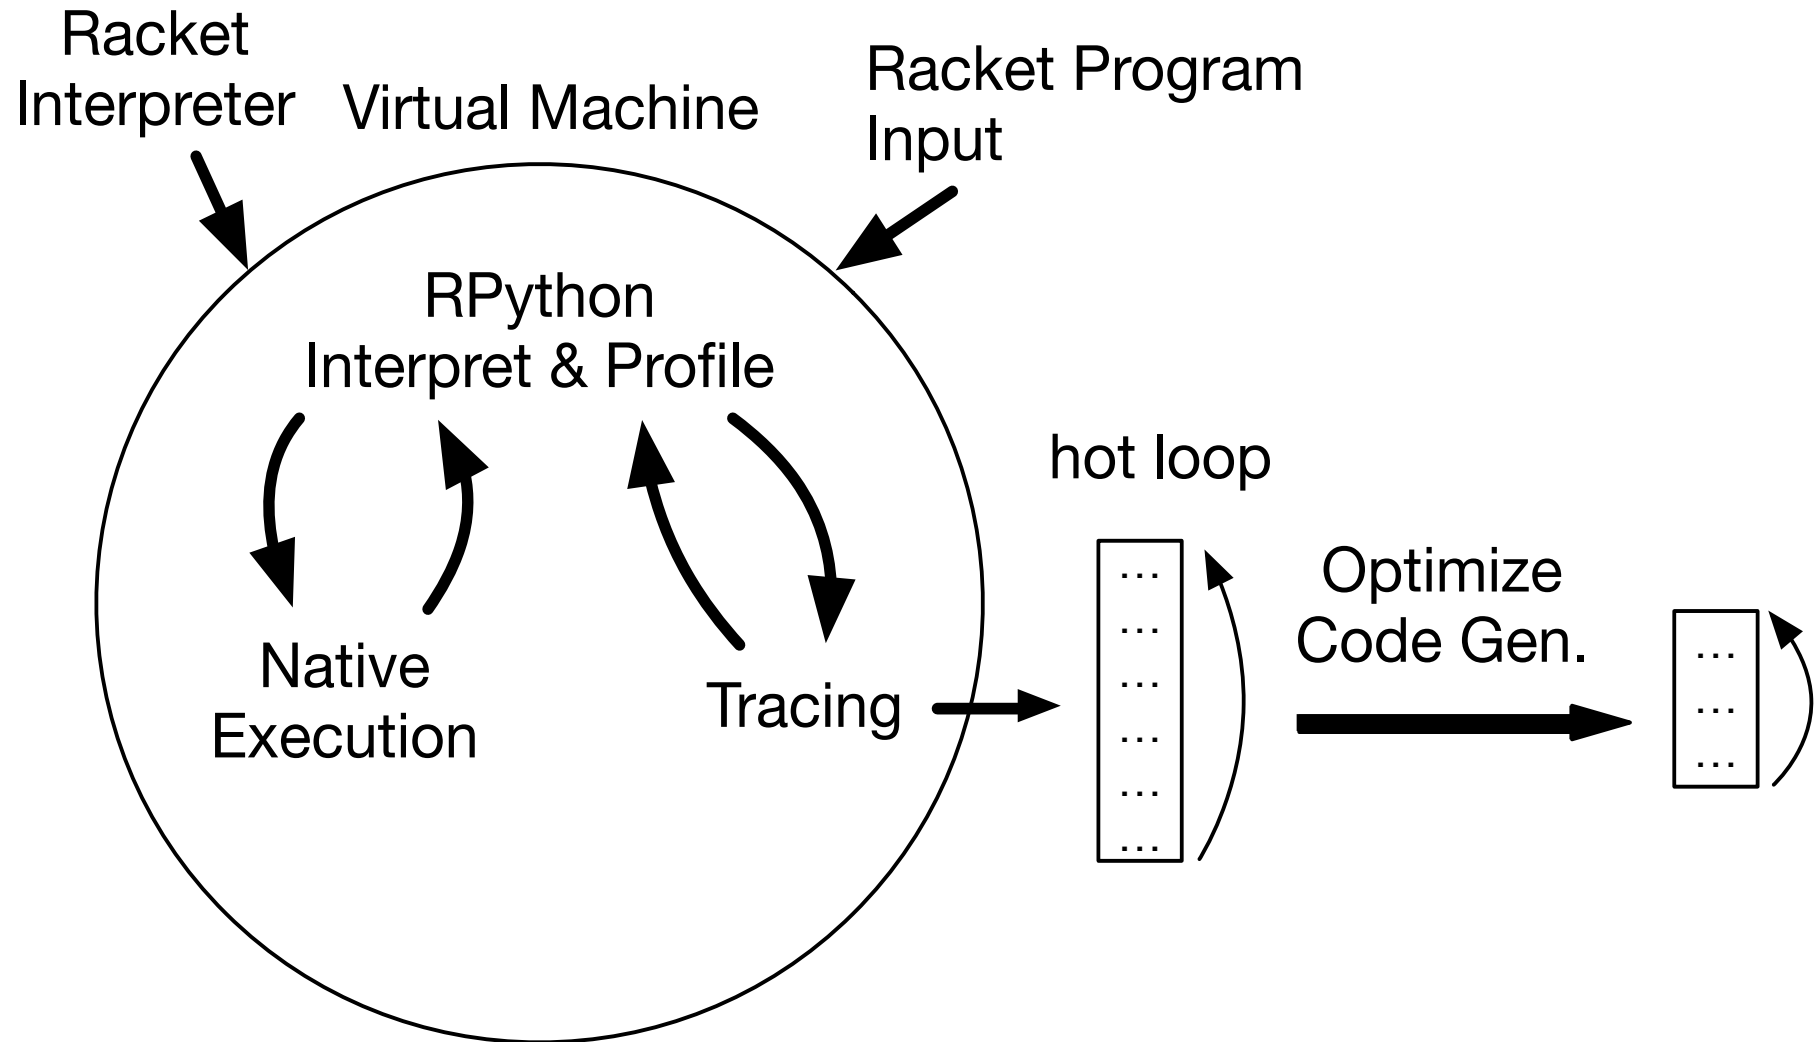
Meta-tracing: the magic part

We didn't write a JIT or an optimizer!

We didn't write a JIT or an optimizer!

RPython creates a JIT from an interpreter

Meta-tracing JIT



CEK Machine

$$e ::= x \mid \lambda x. e \mid e e$$

$$\kappa ::= [] \mid \mathbf{arg}(e, \rho)::\kappa \mid \mathbf{fun}(v, \rho)::\kappa$$

$$\langle x, \rho, \kappa \rangle \longmapsto \langle \rho(x), \rho, \kappa \rangle$$

$$\langle (e_1 e_2), \rho, \kappa \rangle \longmapsto \langle e_1, \rho, \mathbf{arg}(e_2, \rho)::\kappa \rangle$$

$$\langle v, \rho, \mathbf{arg}(e, \rho')::\kappa \rangle \longmapsto \langle e, \rho', \mathbf{fun}(v, \rho)::\kappa \rangle$$

$$\langle v, \rho, \mathbf{fun}(\lambda x. e, \rho')::\kappa \rangle \longmapsto \langle e, \rho'[x \mapsto v], \kappa \rangle$$

CEK Advantages

Fast continuations

Tail recursion

Arbitrary size stack

CEK Advantages

Fast continuations

Tail recursion

Arbitrary size stack

Allocation everywhere

From CEK to JIT

1. Whole-program type inference
2. Translation to C
3. Adding JIT based on hints

Main Interpreter Loop

```
try:
    while True:
        driver.jit_merge_point()
        if isinstance(ast, App):
            prev = ast
            ast, env, cont = ast.interpret(env, cont)
            if ast.should_enter:
                driver.can_enter_jit()
except Done, e:
    return e.values
```

Other hints

Immutable Data

Loop unrolling

Constant functions


Specialization



A loop by any other name

Detecting loops

Record back-edges in control flow graph



The default approach

Detecting loops

Record back-edges in control flow graph



Function calls, not loops

the default approach

Detecting loops

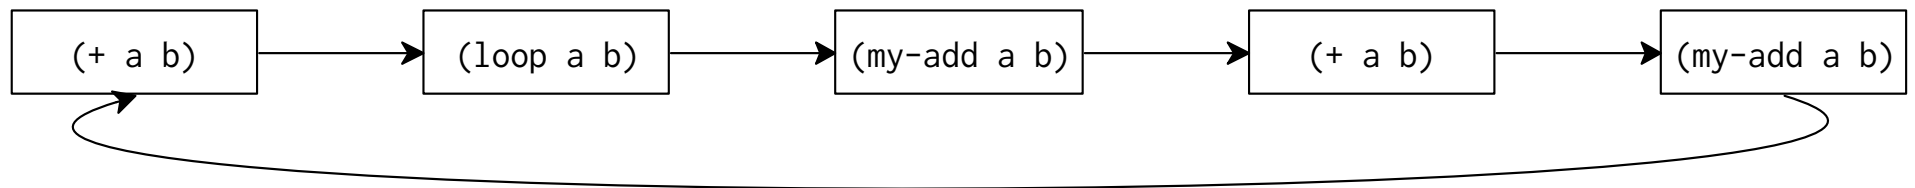
A loop is a repeated AST node

Detecting loops

A loop is a repeated AST node

```
(define (my-add a b) (+ a b))  
(define (loop a b)  
  (if (= a b) 1  
      (loop (my-add a b)  
            (my-add a b)))))
```

Trace from hot node back to itself

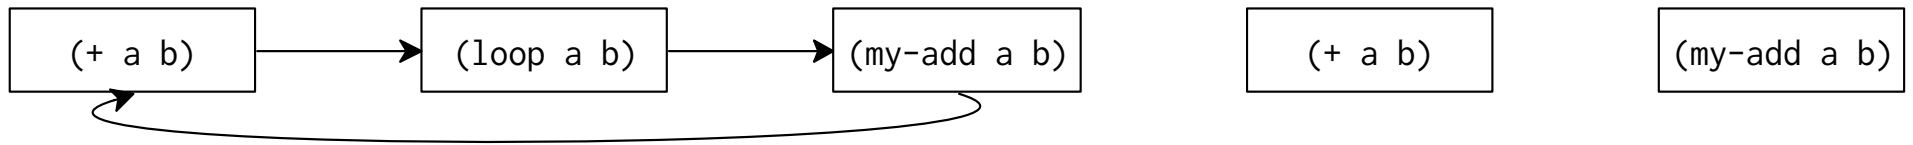


Detecting loops

A loop is a repeated AST node

```
(define (my-add a b) (+ a b))  
(define (loop a b)  
  (if (= a b) 1  
      (loop (my-add a b)  
            (my-add a b)))))
```

Trace from hot node back to itself




Detecting loops

A loop is a repeated AST node


```
(define (my-add a b)
  (define (loop a b)
    (if (= a b)
        (1+ a b)
        (add a b))))
```

Trace not node back to itself

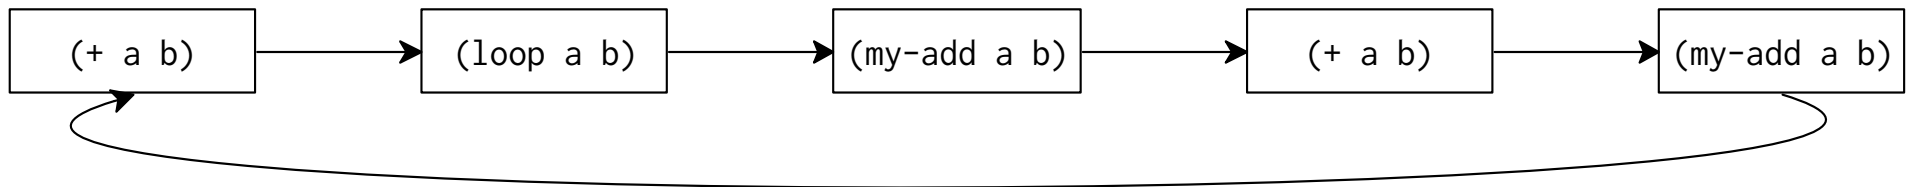


Detecting loops

Construct control flow graph dynamically



Combine with added context





Optimizations

Optimization in the interpreter


A-normalization

Assignment conversion

Environment optimization

Data structure specialization

Storage Strategies



Optimizations we don't do


Closure conversion

Pointer tagging (64-bit integers!)

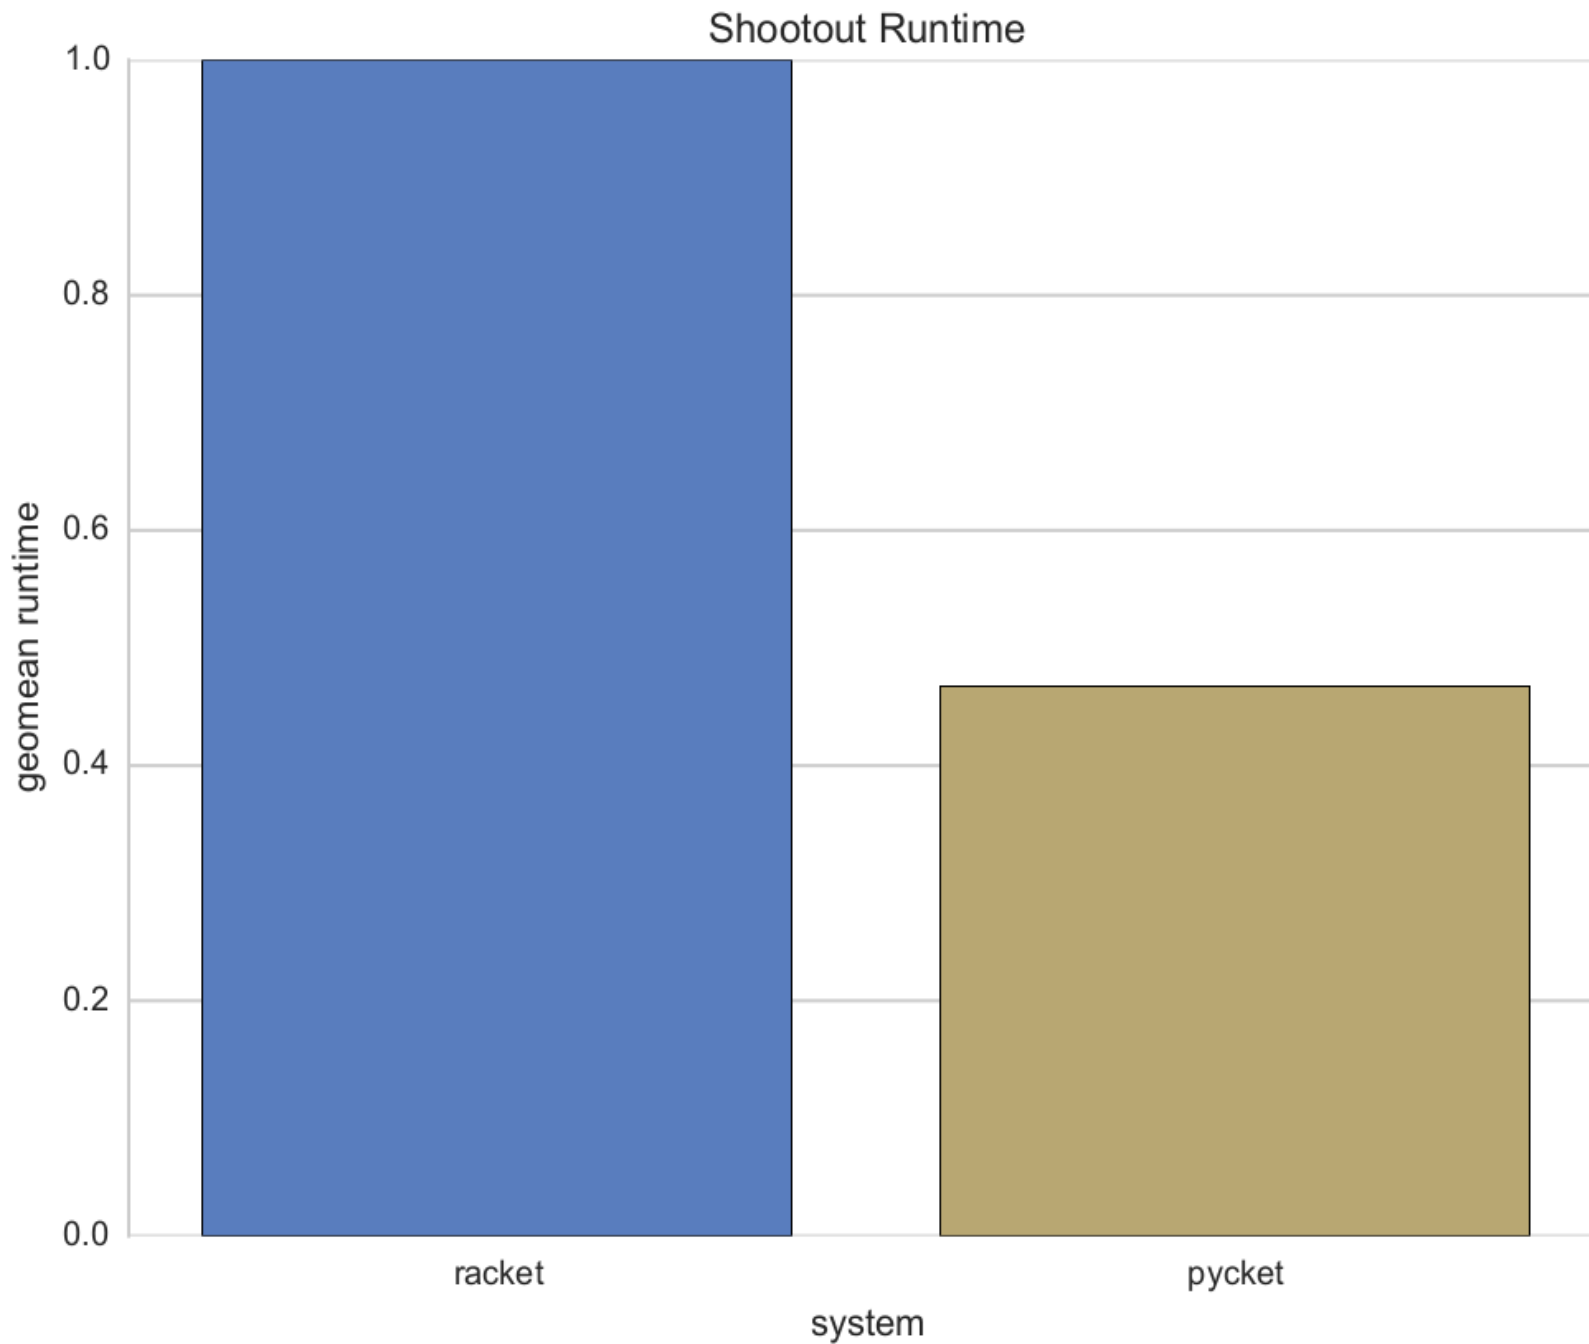


How well does it work?

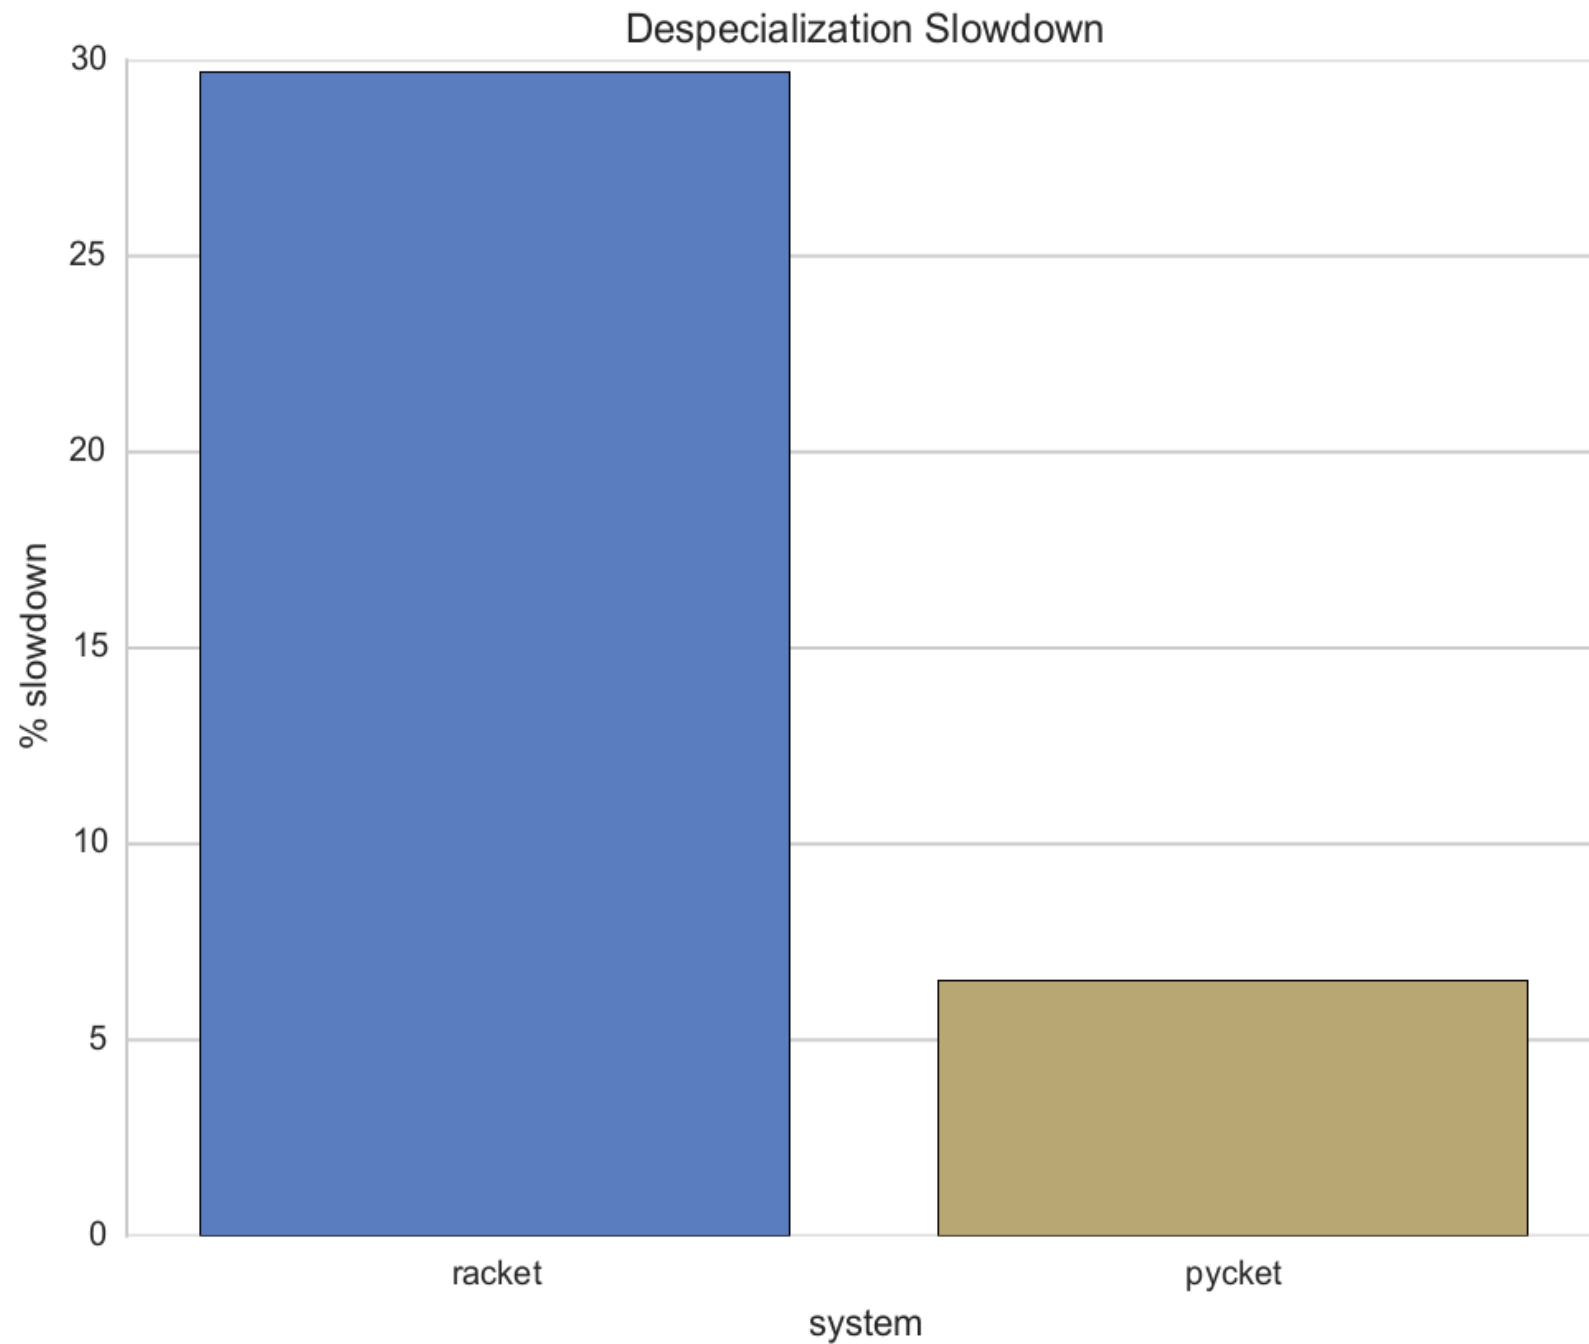
Scheme benchmarks



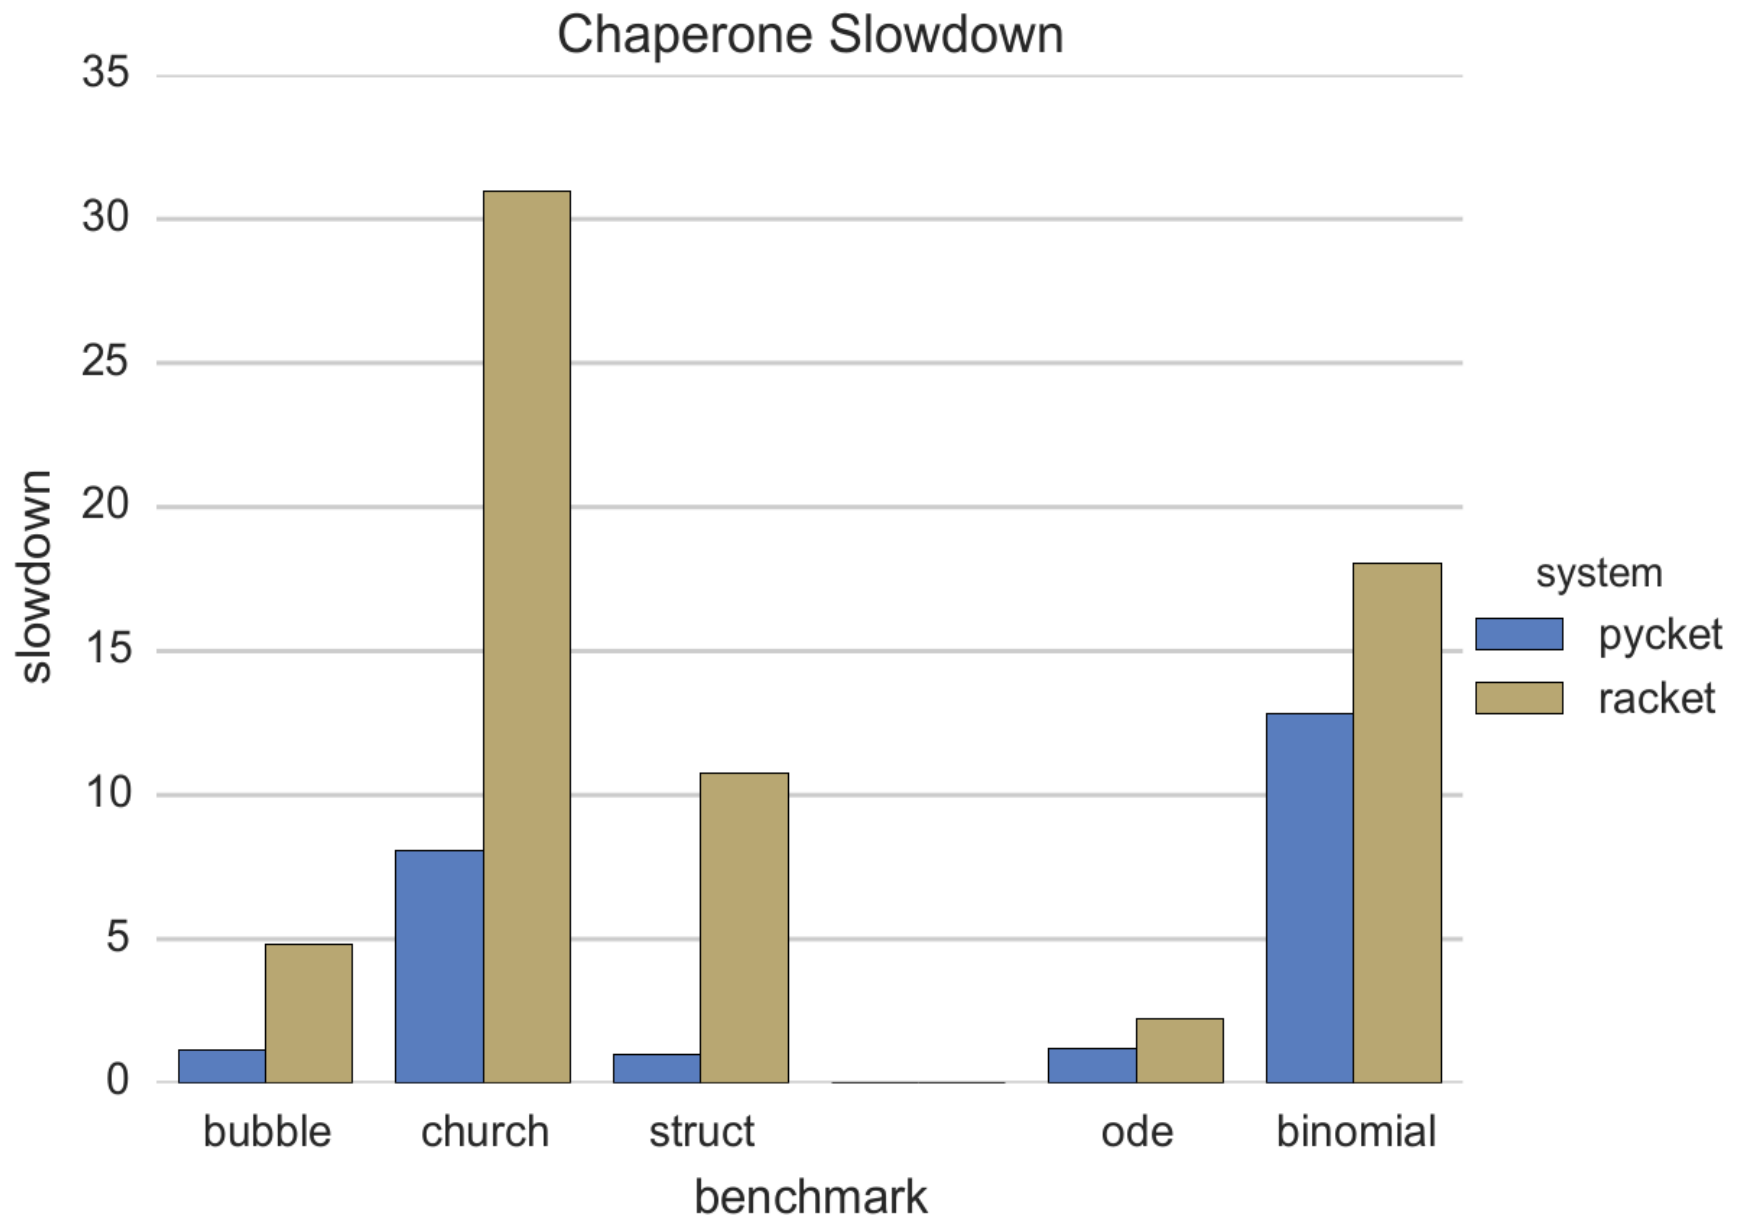
Shootout benchmarks




Shootout benchmarks




Contract benchmarks



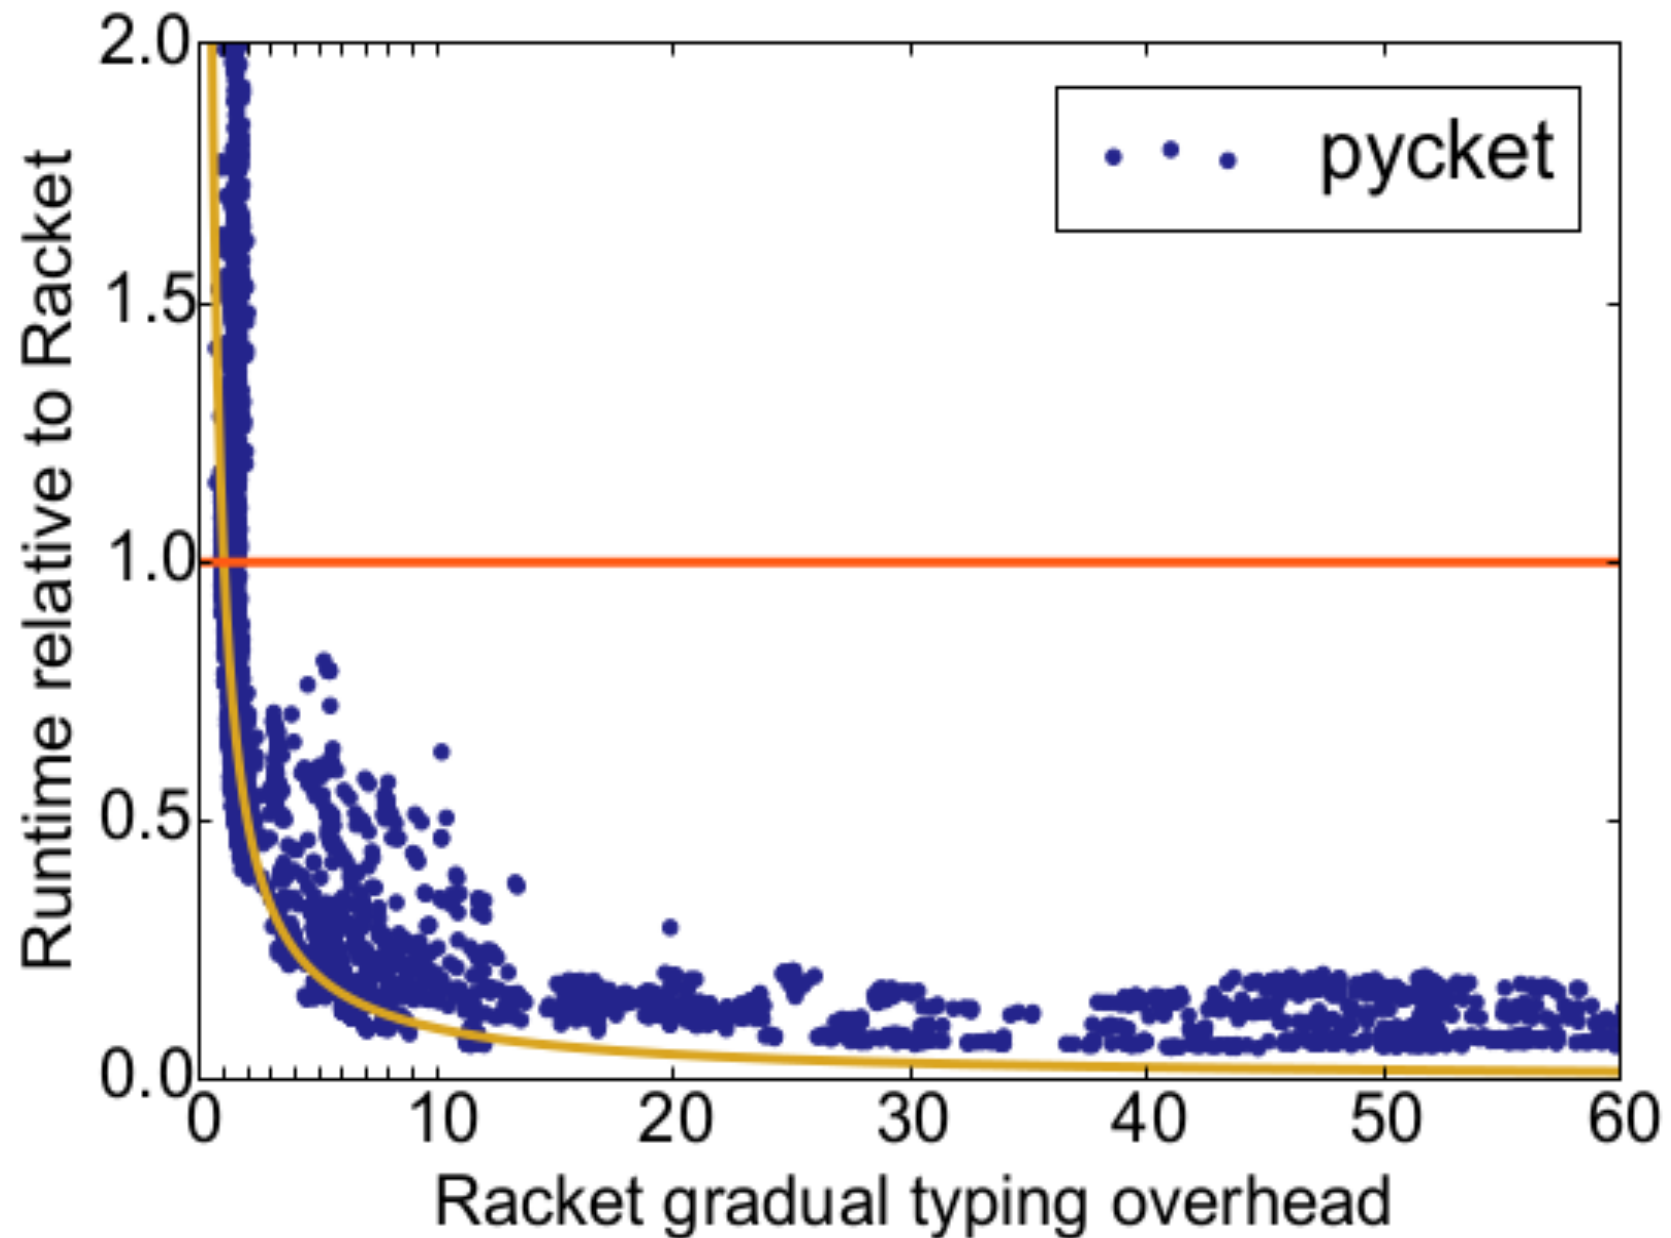
Gradual Typing benchmarks



Gradual Typing benchmarks



Gradual Typing benchmarks





The future of Pycket

What works ...

Basic Scheme

`lambda`, `call-with-values`, `call/cc`, `complex?`

Core Racket

`continuation-marks`, `make-hash`, `contract`

Structures and classes

`struct`, `struct-property`, `object%`, `mixin`

Input/output

`print`, `read`, `call-with-input-file`

Typed Racket

`#lang typed/racket`

Contracts

`chaperone-procedure`, `make-contract`

What doesn't work ...

Concurrency and parallelism
thread, future, place

FFI
make-ctype, editor%

DrRacket

Scribble

Compilation at runtime
eval, compile

Networking
web-server, tcp-connect

Next steps

AOT + JIT = ❤️

Next steps

Expose tracing to programs

Next steps

Accelerate branchy programs

Tracing JIT compilers:

- ★ great for functional languages
- ★ great for generic functions
- ★ great for gradual typing

github.com/samth/pycket